

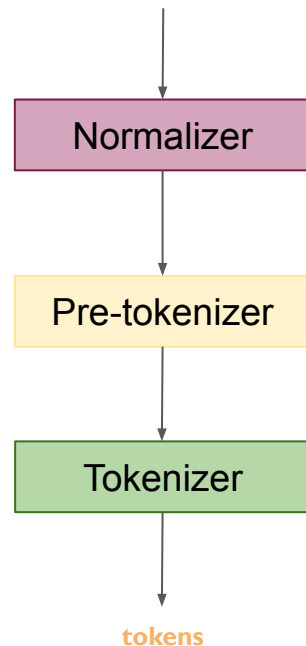
Topics in Natural Language Processing

Look Before You Leap: Pre-tokenization and Tokenization

Dr. Maunendra Sankar Desarkar

Topics

- + Text Normalization
- + Text Tokenization
- + General Tokenization algorithms
 - + Byte-Pair Encoding
 - + Word Piece
 - + Unigram language model
 - + Sentence Piece



Text or Lexical Normalization

Translating non-standard text to a standard form

Why do we need it?

1. Reducing vocabulary size
2. Bring commonality in meaning across multiple inputs where the same word is intended to be used but they appear in different forms

new pix coming tomoroe



new pictures coming tomorrow

Source: [Lexical Normalization | NLP-progress](#) (Ruder)

Datasets for such tasks

1. [LexNorm](#)
2. [MultiLexNorm](#)

Examples from MultiLexNorm

Lang.	Language name	Normalization example
DA	Danish	De skarpe lamper gjorde destromindre ek bedre . De skarpe lamper gjorde destro mindre ikke bedre .
DE	German	ogäj isch hätts auch dwiddern könn Okay ich hätte es auch twittern können
EN	English	u hve to let ppl decide what dey want to do you have to let people decide what they want to do
ES	Spanish	@username cuuxamee sii pero veen yaa eem @username escúchame sí pero ven ya eh
HR	Croatian	svi frendovi mi nešto rade , veceras san osta sam . svi frendovi mi nešto rade , večeras sam ostao sam .
ID-EN	Indonesian-English	pdhal not fully bcs those ppl jg sih . padahal not fully because those people juga sih .
IT	Italian	a Roma è così primavera che sembra gia giov a Roma è così primavera che sembra già giovedì
NL	Dutch	Kga me wss trg rolle vant lachn Ik ga me waarschijnlijk terug rollen van het lachen

Lexical Normalization

Non-standard words (NSWs) are normalised to one or more canonical English words based on a pre-defined lexicon.

- */ o v e* should be normalised to *love* (many-to-one normalisation),
- *tmrw to tomorrow* (one-to-one normalisation),
- *cu* to see you (one-to-many normalisation).
- NOTE: *IBM* should be left untouched as it is in the lexicon and in its canonical form, and the informal *lol* should be expanded to *laughing out loud*.

Original tweet

@USER, r u cuming 2 MidCorner dis Sunday?

Normalized tweet

@USER, are you coming to MidCorner this Sunday?

Original tweet

Still have to get up early 2mr thou 😞 so Gn 🙄

Normalized tweet

Still have to get up early tomorrow though 😞 so Good night 🙄

MoNoise for Lexical Normalization

Two steps - Candidate generation and Ranking

Candidate generation

1. Embedding-based
2. Aspell (edit-distance based)
3. Lookup list from training data
4. Split and check in valid word list
(for longer words)
5. Original word

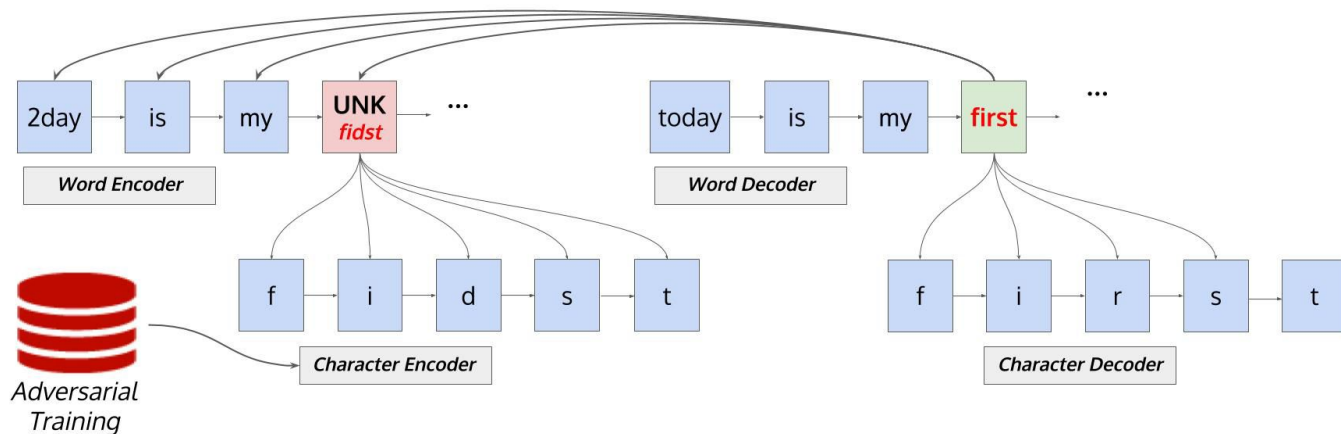
Candidate Ranking

1. Is Original
2. Embedding-distance
3. Aspell Ranking
4. Lookup List
5. N-Gram Probability
6. Character Order Match

What evaluation metrics can be used?

Seq2Seq for Lexical Normalization

- Recent trend - end-to-end
- Can we adapt Seq2Seq Models for this task?
 - Word-based seq2seq: does not work well with OOV/low-probability words
 - Char-based seq2seq: Longer training time. Also, not robust for OOV



More on Lexical Normalization

- Muller, Benjamin, Benoît Sagot, and Djamé Seddah. "[Enhancing BERT for Lexical Normalization](#)." *Proceedings of the 5th Workshop on Noisy User-generated Text (W-NUT 2019)*. 2019.
- Ashmawy, Mohamed, Mohamed Waleed Fakhr, and Fahima A. Maghraby. "[Lexical normalization using generative transformer model \(LN-GTM\)](#)." *International Journal of Computational Intelligence Systems* 16.1 (2023): 183.
- Bikaun, Tyler, Melinda Hodkiewicz, and Wei Liu. "MaintNorm: A corpus and benchmark model for lexical normalisation and masking of industrial maintenance short text." *Proceedings of the Ninth Workshop on Noisy and User-generated Text (W-NUT 2024)*. 2024.
- Bucur, Ana-Maria, Adrian Cosma, and Liviu P. Dinu. "Sequence-to-Sequence Lexical Normalization with Multilingual Transformers." *Proceedings of the Seventh Workshop on Noisy User-generated Text (W-NUT 2021)*. 2021.

Tokenization: Early

In NLP studies, it is conventional to concentrate on pure analysis or generation while taking the basic units, namely words, for granted. It is an obvious truth, however, that without these basic units clearly segregated, it is impossible to carry out any analysis or generation. But too little attention has so far been paid to the process, a kind of preprocessing in a sense, of identifying basic units to be processed. The simplicity of




basic units clearly segregated

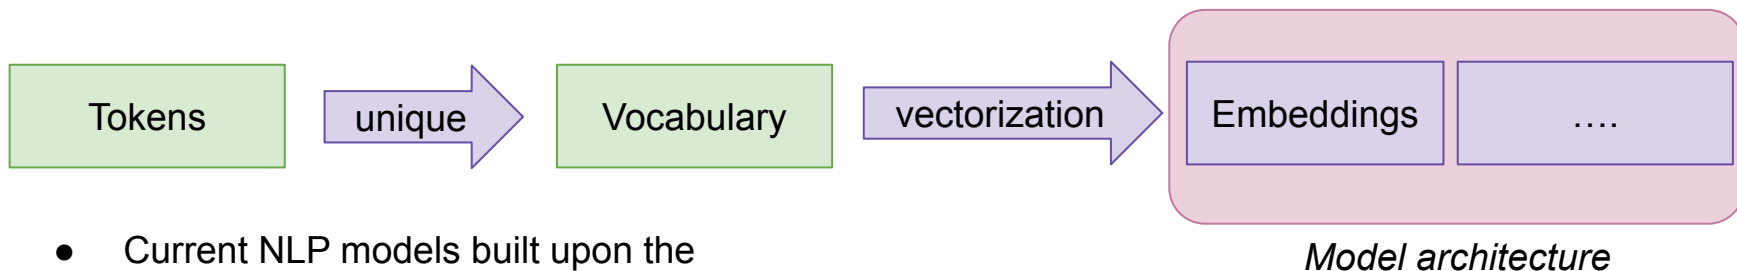
Question: What are the basic units? Or How do we define basic units?

Possible answers:

1. Words (Human way)
2. Characters
3. Subwords
4. Morphemes (Linguistic way)



Tokenization: Big Picture



- Current NLP models built upon the paradigm of:
 - **Pretraining + Fine Tuning**
- Fine Tuning \Rightarrow Typically multiple tasks
- Fixed vocabulary after tokenization (*adaptation techniques exists.*)
- Embeddings of fixed vocabulary are learned

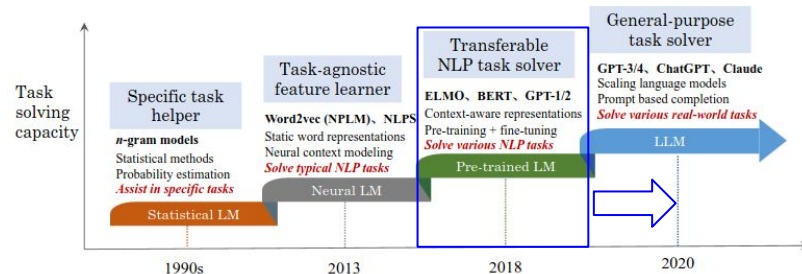


Image credits: <https://arxiv.org/pdf/2303.18223>

Typically in downstream task, same vocabulary as pre training is used \Rightarrow vocabulary should be good as per the downstream task too

Tokenization: What?

Breaking into smaller units or pieces called tokens

or

Identifying *word* boundaries

Ideally: Character tokens would be good 😊

Why?

1. Fixed in numbers
2. Solves Unknown token problem

Cons:

1. Increase in text length
2. Difficulty to capture the word level understanding
3. Existing architectures are not that good

Tokenization: What?

Breaking into smaller units or pieces called tokens

or

Identifying *word* boundaries

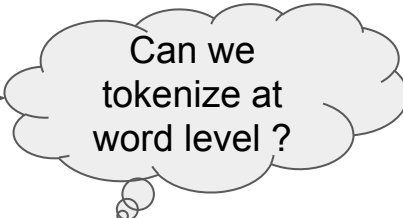
Ideally: Character tokens would be good 😊

Why?

1. Fixed in numbers
2. Solves Unknown token problem

Cons:

1. Increase in text length
2. Difficulty to capture the word level understanding
3. Existing architectures are not that good



Can we
tokenize at
word level ?

Tokenization: What?

Breaking into smaller units or pieces called tokens

or

Identifying *word* boundaries

Ideally: Character tokens would be good 😊

Why?

1. Fixed in numbers
2. Solves Unknown token problem

Cons:

1. Increase in text length
2. Difficulty to capture the word level understanding
3. Existing architectures are not that good

Can we
tokenize at
word level ?

Why? Easy for models to capture semantics

Cons:

1. Large vocabulary
2. Difficult to scale in Multilingual setting (7000+ languages across globe)
3. Unknown token problem exists

Tokenization: What?

Breaking into smaller units or pieces called tokens

or

Identifying *word* boundaries

Ideally: Character tokens would be good 😊

Why?

1. Fixed in numbers
2. Solves Unknown token problem

Cons:

1. Increase in text length
2. Difficulty to capture the word level understanding
3. Existing architectures are not that good

Can we
tokenize at
word level ?

Is there
something
between
Character and
Word level?

Why? Easy for models to capture semantics

Cons:

1. Large vocabulary
2. Difficult to scale in Multilingual setting (7000+ languages across globe)
3. Unknown token problem exists

Tokenization: Subword (the current standard)

- Handles unknown token problem
- Rare words are handled efficiently
- Open vocabulary
- Falls in between Word and Character level
- Controllable vocabulary size
- Variants:
 - Characters
 - Bytes (Currently adapted by LLMs!)

Subword algorithms:

1. Byte-Pair Encoding (BPE)
2. Word Piece
3. Unigram language model
4. Sentence Piece

Tokenization: Subword : BPE

Algorithm 1 Byte-pair encoding (Sennrich et al., 2016; Gage, 1994)

```
1: Input: set of strings  $D$ , target vocab size  $k$ 
2: procedure BPE( $D, k$ )
3:    $V \leftarrow$  all unique characters in  $D$ 
4:     (about 4,000 in English Wikipedia)
5:   while  $|V| < k$  do            $\triangleright$  Merge tokens
6:      $t_L, t_R \leftarrow$  Most frequent bigram in  $D$ 
7:      $t_{\text{NEW}} \leftarrow t_L + t_R$   $\triangleright$  Make new token
8:      $V \leftarrow V + [t_{\text{NEW}}]$ 
9:     Replace each occurrence of  $t_L, t_R$  in
10:       $D$  with  $t_{\text{NEW}}$ 
11:   end while
12:   return  $V$ 
13: end procedure
```

Tokenization: Subword : BPE Example

Steps:


1. Compute unique set of words
2. Build character vocabulary
3. Add new tokens by merging most frequency pairs
4. Continue till the desired vocabulary is achieved

Unique word	Freq.
hug	10
pug	5
pun	12
bun	4
hugs	5



Base characters: unique character in a language.

Here: ["b", "g", "h", "n", "p", "s", "u"] \Rightarrow 7 **character encoded** symbols




h u g	10
p u g	5
p u n	12
b u n	4
h u g s	5

Tokenization: Subword : BPE Example

Vocab: ["b", "g", "h", "n", "p", "s", "u"]

h u g	10
p u g	5
p u n	12
b u n	4
h u g s	5



Most frequent pairs

u g \Rightarrow 20

Merge rule:

"u" "g" \Rightarrow "ug"



New Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug"]

h ug	10
p ug	5
p u n	12
b u n	4
h ug s	5

Most frequent pairs



u n \Rightarrow 16

Steps:

1. Compute unique set of words
2. Build character vocabulary
3. Add new tokens by merging most frequency pairs
4. Continue till the desired vocabulary in achieved

Tokenization: Subword : BPE Example

Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug"]

h ug	10
p ug	5
p u n	12
b u n	4
h ug s	5

Most frequent pairs



u **n** \Rightarrow 16



Updated Merge rule:

"u" "g" \Rightarrow "ug"

"u" "n" \Rightarrow "un"

Steps:

1. Compute unique set of words
2. Build character vocabulary
3. Add new tokens by merging most frequency pairs
4. Continue till the desired vocabulary is achieved

h ug	10
p ug	5
p un	12
b un	4
h ug s	5

New Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug", "un"]

Tokenization: Subword : BPE Example

Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug", "un"]

h ug	10
p ug	5
p un	12
b un	4
h ug s	5

Most frequent pairs



h ug \Rightarrow 15

Updated Merge rule:

"u" "g" \Rightarrow "ug"

"u" "n" \Rightarrow "un"

"h" "ug" \Rightarrow "hug"



hug	10
p ug	5
p un	12
b un	4
hug s	5

Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug", "un", "hug"]

Steps:

1. Compute unique set of words
2. Build character vocabulary
3. Add new tokens by merging most frequency pairs
4. Continue till the desired vocabulary is achieved

Tokenization: Subword : BPE Example

Learned Merge rule:

“u” “g” \Rightarrow “ug”

“u” “n” \Rightarrow “un”

“h” “ug” \Rightarrow “hug”

Vocab: [“b”, “g”, “h”, “n”, “p”, “s”, “u”, “ug”, “un”, “hug”]

How do we tokenize new words?

\Rightarrow Using Merge rules

Examples:

1. bug \Rightarrow [“b”, “ug”]
2. mug \Rightarrow [“UNK”, “ug”]
3. unhug \Rightarrow [“un”, “hug”]

Tokenization: Subword : WordPiece


- Developed by Google
- Used for pretraining popular BERT
- Similar to BPE differs in:
 - Initialization of Base vocabulary
 - Selection of pairs to be merged
 - Encoding process
 - Only saves final vocabulary
 - Uses prefix ## to identify subword

Tokenization: Subword : WordPiece Example

Unique word	Freq.
hug	10
pug	5
pun	12
bun	4
hugs	5



h ##u ##g	10
p ##u ##g	5
p ##u ##n	12
b ##u ##n	4
h ##u ##g ##s	5



Base characters:

- (a) beginning character of a word
- (b) character present inside a word preceded by the ##

Here: ["b", "h", "p", "##g", "##n", "##s", "##u"]

Tokenization: Subword : WordPiece Example

Vocab: ["b", "h", "p", "##g", "##n", "##s", "##u"]

h ##u ##g	10
p ##u ##g	5
p ##u ##n	12
b ##u ##n	4
h ##u ##g ##s	5



Find pairs according to score
& merge

h ##u ##g	10
p ##u ##g	5
p ##u ##n	12
b ##u ##n	4
h ##u ##gs	5

New Vocab: ["b", "h", "p", "##g", "##n", "##s", "##u", "##gs"]

How to compute scores?

$$\text{score}(a,b) = f(a,b) / (f(a) * f(b))$$

$$\begin{aligned}\text{score}("##g", "##s") &= f("##g ##s") / (f("##g") * f("##s")) \\ &= 5/20*5 \\ &= 1/20\end{aligned}$$

Tokenization: Subword : WordPiece Example

Vocab: [“b”, “h”, “p”, “##g”, “##n”, “##s”, “##u”, “##gs”]

$$\text{score}(a,b) = f(a,b) / (f(a) * f(b))$$

h ##u ##g	10
p ##u ##g	5
p ##u ##n	12
b ##u ##n	4
h ##u ##gs	5



Find pairs according to score
& merge

hu ##g	10
p ##u ##g	5
p ##u ##n	12
b ##u ##n	4
hu ##gs	5

New Vocab: [“b”, “h”, “p”, “##g”, “##n”, “##s”, “##u”, “##gs”, “hu”]

$$\begin{aligned}\text{score}(\text{“h”}, \text{“##u”}) &= f(\text{“h ##u”}) / (f(\text{“h”}) * f(\text{“##u”})) \\ &= 15/15*36 \\ &= 1/36\end{aligned}$$

Tokenization: Subword : WordPiece Example

Vocab: [“b”, “h”, “p”, “##g”, “##n”, “##s”, “##u”, “##gs”, “hu”]

hu ##g	10
p ##u ##g	5
p ##u ##n	12
b ##u ##n	4
hu ##gs	5



*Find pairs according to score
& merge*

$$\text{score}(a,b) = f(a,b) / (f(a) * f(b))$$

hug	10
p ##u ##g	5
p ##u ##n	12
b ##u ##n	4
hu ##gs	5

New Vocab: [“b”, “h”, “p”, “##g”, “##n”, “##s”, “##u”, “##gs”, “hu”, “hug”]

$$\begin{aligned}\text{score}(\text{“hu”}, \text{“##g”}) &= f(\text{“hu ##g”}) / (f(\text{“hu”}) * f(\text{“##g”})) \\ &= 10/15 * 10 \\ &= 1/15\end{aligned}$$

$$\begin{aligned}\text{score}(\text{“hu”}, \text{“##gs”}) &= f(\text{“hu ##gs”}) / (f(\text{“hu”}) * f(\text{“##gs”})) \\ &= 5/15 * 5 \\ &= 1/15\end{aligned}$$

Tokenization: Subword : WordPiece Example

Vocab: [“b”, “h”, “p”, “##g”, “##n”, “##s”, “##u”, “##gs”, “hu”, “hug”]

How do we tokenize new words?

⇒ find longest subword in vocab and split

Examples:

1. hugs	⇒	[“hug”, “##s”]	⇒	[“hug”, “##s”]
2. bugs	⇒	[“b”, “##ugs”]	⇒	[“b”, “##u”, “##gs”]
3. mugs	⇒	[UNK]		
4. bum	⇒	[“b”, “##um”]	⇒	[UNK]

Subword Algorithms: Unigram LM

Unigram LM (Kudo, 2018)

Algorithm 2 Unigram LM (Kudo, 2018)

```
1: Input: set of strings  $D$ , target vocab size  $k$ 
2: procedure UNIGRAMLM( $D, k$ )
3:    $V \leftarrow$  all substrings occurring more than
4:     once in  $D$  (not crossing words)
5:   while  $|V| > k$  do            $\triangleright$  Prune tokens
6:     Fit unigram LM  $\theta$  to  $D$ 
7:     for  $t \in V$  do            $\triangleright$  Estimate token 'loss'
8:        $L_t \leftarrow p_\theta(D) - p_{\theta'}(D)$ 
9:       where  $\theta'$  is the LM without token  $t$ 
10:    end for
11:    Remove  $\min(|V| - k, \lfloor \alpha |V| \rfloor)$  of the
12:    tokens  $t$  with highest  $L_t$  from  $V$ ,
13:    where  $\alpha \in [0, 1]$  is a hyperparameter
14:  end while
15:  Fit final unigram LM  $\theta$  to  $D$ 
16:  return  $V, \theta$ 
17: end procedure
```

- Works in opposite direction of BPE & WordPiece
- Starts with large vocabulary and remove symbols until desired vocab size
- At each step train unigram language model from current vocabulary

- Kudo, Taku. "[Subword Regularization: Improving Neural Network Translation Models with Multiple Subword Candidates](#)." ACL 2018.
- Bostrom, Kaj, and Greg Durrett. "[Byte Pair Encoding is Suboptimal for Language Model Pretraining](#)." EMNLP Findings 2020.

Subword Algorithms: Unigram LM : Example

Unique word	Freq.
hug	10
pug	5
pun	12
bun	4
hugs	5

Vocab: [“h”, “u”, “g”, “hu”, “ug”, “p”, “pu”, “n”, “un”, “b”, “bu”, “s”, “hug”, “gs”, “ugs”]

Subword Algorithms: Unigram LM : Example

Vocab: [“h”, “u”, “g”, “hu”, “ug”, “p”, “pu”, “n”, “un”, “b”, “bu”, “s”, “hug”, “gs”, “ugs”]

hug	10
pug	5
pun	12
bun	4
hugs	5

Frequency of subwords:

“h”	⇒	15 (hug, hugs)
“u”	⇒	36
“g”	⇒	20
“hu”	⇒	15
“ug”	⇒	20 (hug, pug, hugs)
“p”	⇒	17
“pu”	⇒	17
“n”	⇒	16
“un”	⇒	16
“b”	⇒	4
“bu”	⇒	4
“s”	⇒	5
“hug”	⇒	15
“gs”	⇒	5
“ugs”	⇒	5 (hugs)

Probability of subword “ug”

= freq. of “ug” / total freq. of subwords

$$= 20 / 210$$

Subword Algorithms: Unigram LM : Example

How probabilities in Unigram are calculated?

Frequency of subwords:

"h"	⇒	15 (hug, hugs)
"u"	⇒	36
"g"	⇒	20
"hu"	⇒	15
"ug"	⇒	20 (hug, pug, hugs)
"p"	⇒	17
"pu"	⇒	17
"n"	⇒	16
"un"	⇒	16
"b"	⇒	4
"bu"	⇒	4
"s"	⇒	5
"hug"	⇒	15
"gs"	⇒	5
"ugs"	⇒	5 (hugs)

Consider word "pug"

Possible tokens: ["p", "u", "g"], ["p", "ug"], ["pu", "g"]

Probabilities:

$$\begin{aligned}\text{prob.}(["p", "u", "g"]) &= \text{prob.}("p") \times \text{prob.}("u") \times \text{prob.}("g") \\ &= (17/120) \times (36/210) \times (20/210) \\ &= 0.0013\end{aligned}$$

$$\begin{aligned}\text{prob.}(["p", "ug"]) &= \text{prob.}("p") \times \text{prob.}("ug") \\ &= (17/210) \times (20/210) \\ &= 0.0077\end{aligned}$$

$$\begin{aligned}\text{prob.}(["pu", "g"]) &= \text{prob.}("pu") \times \text{prob.}("g") \\ &= (17/210) \times (20/210) \\ &= 0.0077\end{aligned}$$

|
Select one with highest
probability
|

hug	10
pug	5
pun	12
bun	4
hugs	5

Subword Algorithms: Unigram LM : Example

At any given stage:

Word	Possible subwords	Unigram Probabilities
hug	["hug"]	→ 0.071428
pug	["pu", "g"]	→ 0.007710
pun	["pu", "n"]	→ 0.006168
bun	["bu", "n"]	→ 0.001451
hugs	["hug", "s"]	→ 0.001701

hug	10
pug	5
pun	12
bun	4
hugs	5

Loss? Negative log likelihood

$$\begin{aligned}\text{Loss} &= 10 \times (-\log(0.071428)) + 5 \times (-\log(0.007710)) + 12 \times (-\log(0.006168)) + 4 \times \\ &\quad (-\log(0.001451)) + 5 \times (-\log(0.001701)) \\ &= 169.8\end{aligned}$$

Subword Algorithms: Unigram LM : Example

Which one to remove? (Exhaustive)

Let's only see two tokens: "pu" or "hug"

hug	10
pug	5
pun	12
bun	4
hugs	5

What happens if we remove "pu"?

Word	Possible subwords	Unigram Probabilities
hug	["hug"] →	0.071428
pug	["p", "ug"] →	0.007710
pun	["p", "un"] →	0.006168
bun	["bu", "n"] →	0.001451
hugs	["hug", "s"] →	0.001701

Loss = 169.8

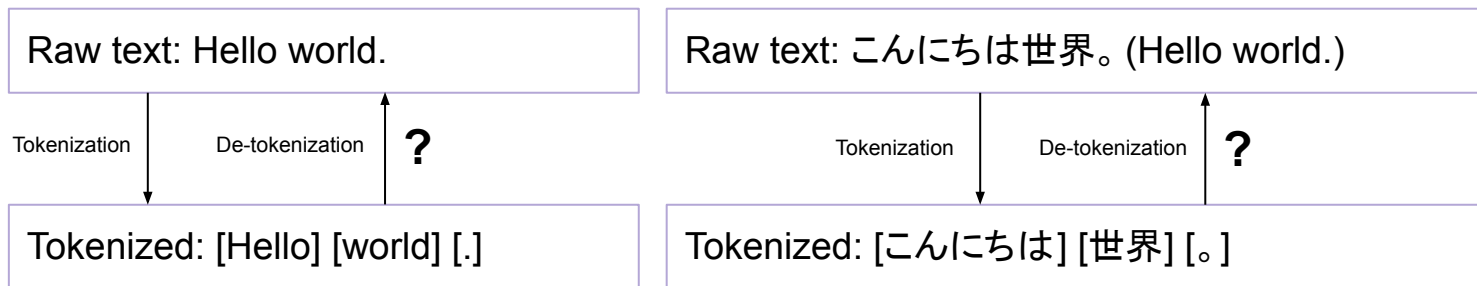
What happens if we remove "hug"?

Word	Possible subwords	Unigram Probabilities
hug	["hu", "g"] →	0.006802
pug	["pu", "g"] →	0.007710
pun	["pu", "n"] →	0.006168
bun	["bu", "n"] →	0.001451
hugs	["hu", "gs"] →	0.001701

Loss = 193.317

Most likely "pu" will be removed

Subword Algorithms: SentencePiece



Observation:

1. Raw text and tokenized sequences are not reversible
2. De-tokenization process are language dependent (Language-specific rules are expensive)

Motivation: How to achieve language-independent lossless de-tokenization?

Subword Algorithms: SentencePiece

DECODE (ENCODE (NORMALIZE (TEXT)) = NORMALIZE (TEXT)

Raw text: Hello world.

Normalize text: Hello_world.

Tokenized: [Hello] [_wor] [ld] [.]

- **NORMALIZE**
 - sequences are treated as a **unicode characters**
 - NFKC-based normalization
 - White Spaces are escaped with _ (Lower One-Eighth block) unicode code point.
- **ENCODE**
 - Uses BPE or Unigram LM algorithm for segmentation
- **DECODE**
 - Substitute _ with space

Unicode codepoint ref.: <https://codepoints.net/U+2581?lang=en>

SentencePiece (Highly recommended): <https://github.com/google/sentencepiece>

Subword Algorithms: Recap

1. BPE (Expansion)
 - a. Starts with **small** base **character** vocabulary
 - b. Merge most frequent pairs
2. WordPiece (Expansion)
 - a. Starts with nearly similarly base byte as BPE
 - b. Merge pairs based on scores
3. Unigram (Reduction)
 - a. Starts with **large** vocabulary
 - b. Remove tokens based on unigram loss

Total vocab. = Base vocab. + New vocab.

Subword Algorithms: Recap

1. BPE (Expansion)
 - a. Starts with **small** base **character** vocabulary
 - b. Merge most frequent pairs
2. WordPiece (Expansion)
 - a. Starts with nearly similarly base byte as BPE
 - b. Merge pairs based on scores
3. Unigram (Reduction)
 - a. Starts with **large** vocabulary
 - b. Remove tokens based on unigram loss

How many characters can a corpus in a specific language have?

Around 4000 in English Wikipedia

Subword Algorithms: Recap

How Small?

1. BPE (Expansion)
 - a. Starts with **small** base **character** vocabulary
 - b. Merge most frequent pairs
2. WordPiece (Expansion)
 - a. Starts with nearly similarly base byte as BPE
 - b. Merge pairs based on scores
3. Unigram (Reduction)
 - a. Starts with **large** vocabulary
 - b. Remove tokens based on unigram loss

How many characters can a corpus in a specific language have?

Around 4000 in English Wikipedia
(Very Large)

What about characters in other languages?

Subword Algorithms: Recap

How Small?

1. BPE (Expansion)
 - a. Starts with **small** base **character** vocabulary
 - b. Merge most frequent pairs
2. WordPiece (Expansion)
 - a. Starts with nearly similarly base byte as BPE
 - b. Merge pairs based on scores
3. Unigram (Reduction)
 - a. Starts with **large** vocabulary
 - b. Remove tokens based on unigram loss

What about characters in other languages?

149K unicode code point (v15.1)
(**Very large**)




As of [Unicode](#) version 15.1, there are 149,878 [characters](#) with [code points](#), covering 161 modern and historical [scripts](#), as well as multiple symbol sets. This article includes the 1,062 characters in the Multilingual European Character Set 2 ([MES-2](#)) subset, and some additional related characters.

Subword Algorithms: Recap

1. BPE (Expansion)
 - a. Starts with **small** base **Byte** vocabulary
 - b. Merge most frequent pairs
2. WordPiece (Expansion)
 - a. Starts with nearly similarly base byte as BPE
 - b. Merge pairs based on scores
3. Unigram (Reduction)
 - a. Starts with **large** vocabulary
 - b. Remove tokens based on unigram loss

149K unicode code points

Very large set to be considered as base vocabulary for BPE



Solution: Unicode
Encoding

Represent in **Bytes** instead of
Characters.

UTF-8, UTF-16, UTF-32

**UTF-8 can represent all unicode
code points in 1-4 Bytes**

ByT5

